

Data Scraping

Scraping: download_data.py

Libraries: urllib2, BeautifulSoup, Selenium

usage: download_data.py [-h] [-o dir] -d {MX,TW,PE,Rio,PAHO,TN, CME}

optional arguments:

-h, --help show this help message and exit
-o dir, --output dir the output directory
-d {MX,TW,PE,Rio,PAHO,TN,CME}, --dataset {MX,TW,PE,Rio,PAHO,TN,CME}
the data set to download

Example:

download_data.py -d PAHO -o paho/PDFs/

Converting: PDF2CSV.jar

Libraries: PDFBox, Tabula-java, OpenCV

usage: tabula [-b <DIRECTORY>] [-d] [-f <FORMAT>] [-g] [-h] [-i] [-l] [-o
<OUTFILE>] [-p <PAGES>] [-s <PASSWORD>] [-t] [-u] [-v]

Tabula helps you extract tables from PDFs

-b,--batch <DIRECTORY>	Convert all .pdfs in the provided directory.
-d,--debug	Print detected table areas instead of processing.
-f,--format <FORMAT>	Output format: (CSV,TSV,JSON). Default: CSV
-g,--guess	Guess the portion of the page to analyze per page.
-h,--help	Print this help text.
-i,--silent	Suppress all stderr output.
-l,--lattice	Force PDF to be extracted using lattice-mode extraction (if there are ruling lines separating each cell, as in a PDF of an Excel spreadsheet)
-o,--outfile <OUTFILE>	Write output to <file> instead of STDOUT. Default: -
-p,--pages <PAGES>	Comma separated list of ranges, or all. Examples: --pages 1-3,5-7, --pages 3 or --pages all. Default is --pages 1
-s,--password <PASSWORD>	Password to decrypt document. Default is empty
-t,--stream	Force PDF to be extracted using stream-mode extraction (if there are no ruling lines separating each cell)
-u,--use-line-returns	Use embedded line returns in cells. (Only in spreadsheet mode.)
-v,--version	Print version and exit.

Summary: Tabula is a PDF table extractor. Their library tabula-java uses PDFBox. They have a command line application, but I found it lacking in the automatic table detection algorithm. Their algorithm used overly simplistic image processing methods in order to detect vertical and horizontal lines on the page. So in PDF2CSV.jar I use OpenCV to accomplish the line detection much more effectively. I copied most of the code from Tabula-java's command line app, changing the image processing methods and eliminating a few unneeded command line options.

Tabula-java has an '-a,-area' option which allows the user to indicate by coordinates the region of interest on the page. But the whole reason I wrote this program was to automate that process. So that option has been removed. It has a similar '-c,-column' option which I removed for the same reason. I removed two deprecated options '-r' and '-n' which have been replaced with '-l,-lattice' and '-t,-stream'. Option '-g,-guess' should always be used, as that is where the table detection algorithm is called.

Comparison: This is the most general, fast, and accurate of the three conversion programs I have made.

- PDF2CSV.py is very slow (but okay on small tables)
- The first version of PDF2CSV.jar (with no Tabula) is not nearly general enough and would require lots more work (which Tabula has already done!).

Small adjustments needed:

- Handle certain fonts
- Detect inversely colored tables. (Right now it only detects dark lines).

Converting: HTML2CSV.py

Libraries: HTMLParser

usage: HTML2CSV.py [-h] [-o dir] [-r] file [file ...]

A coarse HTML tables to CSV (Comma-Separated Values) converter.

positional arguments:

file the HTML file(s) you want to convert to CSV.

optional arguments:

-h, --help show this help message and exit
 -o dir, --output dir the output directory
 -r, --recursive To get all leaf tables from HTML with nested tables

Examples : HTML2CSV.py -i mypage.html
 : HTML2CSV.py -i *.html -o CSVs/

Automating: crontab

Crontab is a list of commands that you want to run on a regular schedule.

Crontab:

- use 'env EDITOR=vim crontab -e' to add/change a scheduled script
 - use 'crontab -l' to display all scheduled events

```
*      *      *      *      *      command to be executed
-      -      -      -      -
|      |      |      |      |
|      |      |      |      +----- day of week (0 - 6) (Sunday=0)
|      |      |      +----- month (1 - 12)
|      |      +----- day of month (1 - 31)
|      +----- hour (0 - 23)
+----- min (0 - 59)
```

Example:

Including this line in your crontab will run command every day at noon:

```
0 12 * * * cd ~/Dengue/Peru/Code/Python && ./download_PDFs_weekly.py
```

Converting: PDF2CSV.py (depricated in favor of Java)

Libraries: pdfminer

usage: PDF2CSV.py [-h] [-o dir] [-s] (-b | -c) file [file ...]

positional arguments:

file the input file(s)

optional arguments:


-h, --help show this help message and exit
-o dir, --output dir the output directory
-s, --separate Each page of PDF is a separate table
-b, --boundary differentiate cells by cell borders
-c, --coordinate differentiate cells by text element coordinates

Examples : PDF2CSV.py country/PDFs/* -o country/CSVs/ -b
 : PDF2CSV.py country/PDFs/weekly_pages.pdf -o country/weekly-CSVs/ -c -s

This script uses the pdfminer python library. Pdfminer extracts the contents of PDFs. From (pg. 12-15): <https://media.readthedocs.org/pdf/pdfminer-docs/latest/pdfminer-docs.pdf>

A layout analyzer returns a LTPage object for each page in the PDF document. This object contains child objects within the page, forming a tree structure.

pdfminer LT objects:



LTPage

- Represents an entire page. May contain child objects like LTextBox, LFigure, LImage, LRect, LCurve and LLine.

LTextBox

- Represents a group of text chunks that can be contained in a rectangular area. Note that this box is created by geometric analysis and does not necessarily represents a logical boundary of the text. It contains a list of LTextLine objects. get_text() method returns the text content.

LTextLine

- Contains a list of LChar objects that represent a single text line. The characters are aligned either horizontally or vertically, depending on the text's writing mode. get_text() method returns the text content.

LChar

- Represents a single character of text.

LTAnno

- Represent an actual letter in the text as a Unicode string. Note that, while a LChar object has actual boundaries, LTAnno objects does not, as these are “virtual” characters, inserted by a layout analyzer according to the relationship between two characters (e.g. a space).

LTFigure

- Represents an area used by PDF Form objects. PDF Forms can be used to present figures or pictures by embedding yet another PDF document within a page. Note that LTFigure objects can appear recursively

LTImage

- Represents an image object. Embedded images can be in JPEG or other formats, but currently PDFMiner does not pay much attention to graphical objects.

LTLine

- Represents a single straight line. Could be used for separating text or figures.

LTRect

- Represents a rectangle. Could be used for framing another pictures or figures.

LTCurve

- Represents a generic Bezier curve

LT-elements all share some of these public attributes:

- bbox: (x0, y0, x1, y1) <- coordinates of the element's bounding box.
- width
- height
- text (or get_text()) <- for character and text elements only.
- ...

Implementation

PDF2CSV.py has two methods for extracting tabular structures from PDFs.

1. By Cell Boundaries [-b]

This method uses LTLine(s) and LTRect(s) along with their bbox, width, and height attributes to determine cell boundaries. It works well if there are lines/rectangles that divide every cell. See this tutorial for an explanation.

2. By Cell Coordinates [-c]

This method uses LTTextLine(s) and LTChar(s) along with their bbox attribute to determine cell positions. Basically each LTTextLine is sorted first by its y-coordinate (i.e. by row), and then by its x-coordinate (i.e. by column). The text attribute is then extracted from this 2-dimensional, ordered list.

To convert this README to PDF, use: `$ pandoc README.md -o README.pdf -V geometry:margin=0.8in`